

WEEK #1

- 1) What is ML.
- 2) Introductn
- 3) Quiz #1
- 4) Other materials
- 5) Model and Cost Fxn
- 6) Paramtr Learning
- 7) Review Quiz
- 8) Linear Algebra review
- 9) Review Quiz

WEEK #4

- 1) Motivatns
- 2) NN
- 3) Applicatns
- 4) Review (Quiz+PA)

WEEK #7

- 1) Large Margin Class'n.
- 2) Kernels
- 3) SVM's in Practice
- 4) Review (QUIZ+P.A.)

WEEK #10

- 1) Grad. Desc. with Large Datasets
- 2) Advanced Topics
- 3) Review (QUIZ)

WEEK #2

- 1) Environment Setup Instr.
- 2) Multivariate Linear Regr.
- 3) Computing Paramtrs Analyt.
- 4) Submitting Progr. Assignts.
- 5) Review Quiz
- 6) Octave/MATLAB Tutorial
- 7) Review Quiz
- Programming Assignment

WEEK #5

- 1) Cost fcn & Backpropagatn
- 2) Backpropagatn in practice
- 3) Applicatn of NN
- 4) Review (QUIZ + P.A.)

WEEK #8

- 1) Clustering
- 2) Review (QUIZ)
- 3) Motivatn
- 4) PCA Analysis
- 5) Applying PCA
- 6) Review (QUIZ+P.A.)

WEEK #11

- 1) Photo OCR
- 2) Review (QUIZ)
- 3) Conclusion.

WEEK #3

- 1) Classificatn & Representn.
- 2) Logistic Regression Model
- 3) Multiclass Classificatn
- 4) Review Quiz
- 5) Solving the Prob. Over fitting
- 6) Review Quiz
- Programming Assignment.

WEEK #6

- 1) Evaluating a Learn. Algor.
- 2) Bias vs. Variance
- 3) Review (QUIZ+P.A.)
- 4) Building a Spam Classifier
- 5) Handling Skewed Data
- 6) Using Large Data Sets
- 7) QUIZ

WEEK #9

- 1) Density Estimatin
- 2) Building a Anomaly Det.S.
- 3) Multivariate Gauss. Dist.
- 4) Review (QUIZ)
- 5) Predicting Movie Ratings
- 6) Collaborate Filtering
- 7) Low Rank Mtx Factoriztn
- 8) Review (QUIZ+P.A.)

①. What is Machine Learning


- Science of getting computers to learn without being explicitly programmed.
- NN mimics the human brain.

②. INTRODUCTION

2.1. Welcome

- Web search engine: Google, Bing
- Facebook, Apple face recognitn
- Spam filters
- Autonomous robots
- Computational biology
- Handwriting recognitn.
- NLP: Natural Language Processing
- Computer Vision
- Self-customizing programs
- In \forall fields of engineering there is a huge amount of data sets, that we are trying to understand using learning algorithms.


2.2. Supervised Learning.



ex/ Breast cancer
(malignant, benign)

- \Rightarrow Supervised Learning: "right answers" given
- \Rightarrow Regression: Predict cont's valued outputs
- \Rightarrow Classification: Discrete valued outputs (0 or 1) / (0, 1, 2, 3)
- \Rightarrow # of features may vary by problem

2.3. Unsupervised Learning:



ex/ Google News, Social network analysis,
Astronomical data analysis, Market segmentatn.

- \Rightarrow Unsupervised Learning: we are given a data that does not have any labels/ that \forall has the same label.

Cocktail Party Problem Algorithm:

$$[W, s, v] = \text{svd}((\text{repmat}(\text{sum}(X, *X, 1), \text{size}(X, 1), 1) .* X) * X');$$

③. QUIZ #1 (100% ✓)

④. OTHER MATERIALS:

The Course Wiki


5. MODEL & COST FUNCTION

5.1. Model Representatn:

m = # of training examples

x 's = "input" variable/features

y 's = "output" / "target" variable.



$$h_{\theta}(x) = \theta_0 + \theta_1 x \text{ (univariate linear regression)}$$

5.2. Cost Functn:

θ_0, θ_1 : parameters.

→ Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for our training examples

$$\text{minimize}_{\theta_0, \theta_1} \sum_{i=1}^M (h_{\theta}(x^{(i)}) - y^{(i)})^2 \cdot \frac{1}{2m} = J(\theta_0, \theta_1)$$

← Cost func.

(squared error func'n)


5.3. Cost Funct'n - Intuitn1:

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$


Params: θ_0, θ_1

$$\text{Cost Fxn: } J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1



5.4. Cost Funct'n - Intuitn2:



6. PARAMETER LEARNING

6.1. Gradient Descent:

- for minimizing some arbitrary fcn J . ✓ Have some fcn $J(\theta_0, \theta_1)$ ✓ Want $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

Outline:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$ until we hopefully end up at a minimum.

⇒ If you are standing at the pit on the hill, look around and find that the best dirn is to take a little step downhill is roughly that dirn.

Gradient Descent Algorithm:

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j=0 \& j=1)$$

}

learning rate (step size)

⇒ simultaneous update:


$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$

6.2. Gradient descent Intuition:



if derivative ≥ 0

$$\theta_1 := \theta_1 - \alpha (\text{pos'ive } \#) \rightarrow \text{decrease } \theta_1$$

if derivative < 0

$$\theta_1 := \theta_1 - \alpha (\text{neg'ive } \#) \rightarrow \text{increase } \theta_1$$

⇒ if α is too small, gradient descent can be slow

⇒ if α is too large, gradient descent can overshoot the minimum. It may fail to converge or even diverge.

⇒ Gradient descent can converge to a local minimum, even with the learning rate α fixed

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

As we approach a local minimum, gradient descent will automatically take smaller steps. So, no need to decrease α over time.

6.3. Gradient descent for Linear Regression:

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \cdot \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{\partial}{\partial \theta_j} \cdot \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

$$\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \quad \underline{j=0}$$

$$\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \quad \underline{j=1}$$

"Batch" Gradient descent: Each step of gradient descent uses θ the training examples.

⇒ Normal eq'n method can solve this minimization problem numerically, but Gradient descent can scale better to larger data sets than that normal eq'n method.