

# ① EVALUATING A LEARNING ALGORITHM.

## 1.1. Deciding What to Try Next.

→ If you are developing ML, stms, that you know how to choose one of the most promising avenues to spend your time pursuing.

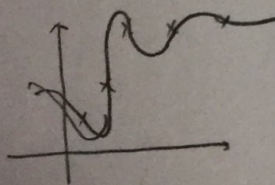
⇒ Debugging a learning algorithm.

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^m \theta_j^2 \right]$$

- Get more training examples
- Try smaller sets of features (to prevent overfitting)
- Try getting additional features (may be current features aren't informative enough & you want to collect more data in the sense of getting more features).
- Try adding polynomial features ( $x_1^2, x_2^2, x_1 x_2$ , etc.)
- Try decreasing  $\lambda$
- Try increasing  $\lambda$

⇒ M.L. diagnostic: test that you can run to gain insight what is/isn't working with a learning algorithm, and gain guidance as to how best to improve its performance.

## 1.2. Evaluating a Hypothesis:



Fails to generalize to new examples not in training set  
How to be sure about overfitting?

Divide the training set into 2 groups: → training set (70%)  
→ test set (30%)

- Learn paramtr  $\theta$  from training data
- Compute test error:

## 1.3. Model Selectn and Train/Validatn/Test Sets:

⇒ Model Selection:

$d$  = degree of polynomial.

$d=1$  → Training error ( $\theta^{(1)}$ ) → Test error ( $\theta^{(1)}$ )

$d=5$  → Training error ( $\theta^{(5)}$ ) → Test error ( $\theta^{(5)}$ )

$J_{\text{test}}(\theta^{(5)})$

→  $J_{\text{test}}(\theta^{(5)})$  is like to be an optimistic estimate of generalizatr error, i.e. our extra paramtr ( $d$ ) is fit to test set.



→ Divide the dataset into 3 groups: training set (60%) →  $J_{train}(\theta)$   
 cross-validation set (20%) →  $J_{cv}(\theta)$   
 test set (20%) →  $J_{test}(\theta)$

Use CV set to select the d-paramtr (rather than using test set).

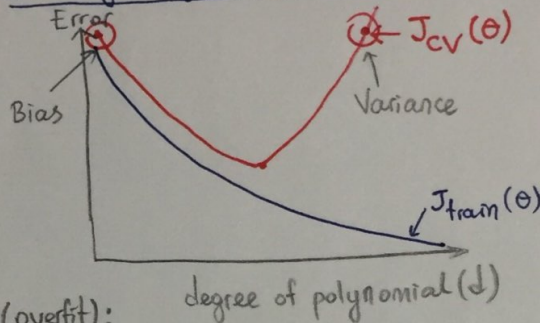
Pick the optimum hypothesis ( $ex/ \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$ )

Estimate generalizatr error for test set  $J_{test}(\theta^{(4)})$

## ② BIAS Vs. VARIANCE

### 2.1. Diagnosing bias Vs. variance:

bias → underfitting ( $d=1$ )  
 variance → overfitting ( $d=10$ )  
 just right ( $d=5$ )



Bias (underfit):

$J_{train}(\theta) \rightarrow$  high  
 $J_{cv}(\theta) \rightarrow$  high

Variance (overfit):

$J_{train}(\theta) \rightarrow$  low  
 $J_{cv}(\theta) \gg J_{train}(\theta)$

### 2.2. Regularizatr and Bias/Variance:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Large  $\lambda$ :

High bias, underfit  
 $\lambda = 10^6, \theta_1 \approx 0, \theta_2 \approx 0$   
 $h_{\theta}(x) \approx 0$

Intermediate  $\lambda$ :

Just Right

Small  $\lambda$ :

High Variance, Overfit  
 $\lambda = 0$

⇒ Choosing the regularizatr error:

- 1)  $\lambda = 0 \rightarrow \min_{\theta} J(\theta) \rightarrow \theta^{(1)} \rightarrow J_{cv}(\theta^{(1)})$
- 2)  $\lambda = 0.01$
- 3)  $\lambda = 0.02$
- 4)  $\lambda = 0.04 \rightarrow \min_{\theta} J(\theta) \rightarrow \theta^{(4)} \rightarrow J_{cv}(\theta^{(4)})$
- 2)  $\lambda = 10 \rightarrow \theta^{(12)} \rightarrow J_{cv}(\theta^{(12)})$

Pick (say)  $\theta^{(5)} \rightarrow$  Test Error:  $J_{test}(\theta^{(5)})$

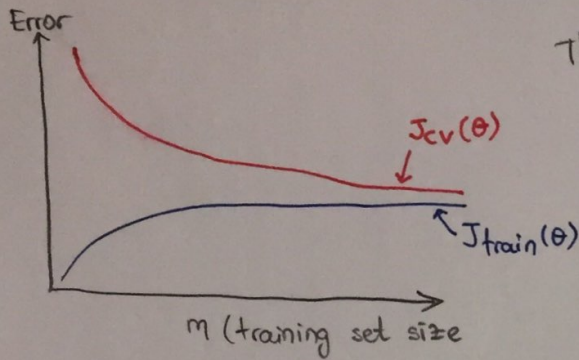




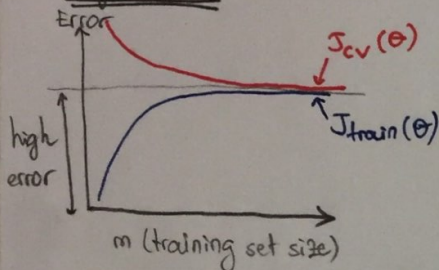
### 2.3. Learning Curves:

25

The more data you have  $\Rightarrow$  the better the hypothesis you fit.



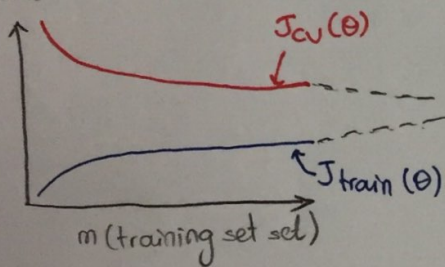
#### High Bias:



$J_{cv}$  &  $J_{train}$  will be very similar

If a learning algorithm is suffering from high bias, getting more training data will not (by itself) help much.

#### High Variance:



$\rightarrow$  Large gap btw training error & cv-error

If a learning algorithm is suffering from high variance, getting more training is likely to help.

### 2.4. Deciding What to do Next (Revisited)

- $\rightarrow$  Get more training examples  $\rightarrow$  fixes high variance
- $\rightarrow$  Try smaller sets of features  $\rightarrow$  fixes high variance
- $\rightarrow$  Try getting additional features  $\rightarrow$  fixes high bias
- $\rightarrow$  Try adding polynomial features  $\rightarrow$  fixes high bias
- $\rightarrow$  Try decreasing  $\lambda \rightarrow$  fixes high bias
- $\rightarrow$  Try increasing  $\lambda \rightarrow$  fixes high variance

#### Small NN

- computationally cheaper
- prone to underfitting

#### Large NN

- more expensive
- prone to overfitting
- Use regularization  $\nearrow$
- Select # of hidden layers.

③ REVIEW: QUIZ + P.A.



# ④ BUILDING A SPAM CLASSIFIER.

26

## 4.1. Prioritizing What to Work On:

1)  $x$  = features of email,  $y = 1$  (spam) / 0 (non-spam)

features  $x$ : choose 100 words indicative of spam/non-spam.

$$x = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \\ 1 \\ 0 \\ \vdots \end{bmatrix}$$

2) collect lots of data

3) develop sophisticated features based on email routing info.

4) " " " for message body ("discounts" ...)

5) " " algorithm to detect misspellings (medicine)

## 4.2. Error Analysis:

1) start with simple algorithm

2) plot learning curves

3) error analysis: manually examine the examples

ex/ 100 errors on categorization of spam classifier:

i) what type of mail it is

ii) what cues (features) you think would have helped the algorithm to classify.

## 4.3. Error metrics for skewed classes:

Find that you got 1% error on test set.

→ Only 0.50% of patients have cancer.

↳ skewed classes.

Function  $y = \text{predictCancer}(x)$

end  $y = 0$ ; → 0.5% error gives

recall = 0

Precision / Recall. ( $y = 1$  in presence of rare class).

$y = 1$  in presence of rare class that we want to detect.

		Actual class	
		1	0
Predicted Class	1	True positive	False positive
	0	False negative	True negative

Precision:

↑ good

$$\frac{\text{true positive}}{\# \text{ predicted positive}} = \frac{\text{True positive}}{\text{True pos} + \text{False pos}}$$

Recall:

↑ good

$$\frac{\text{True pos}}{\# \text{ actual pos}} = \frac{\text{True pos}}{\text{True pos} + \text{False neg.}}$$



### 5.2. Trading Off Precision and Recall.

Logistic regression:  $0 \leq h_{\theta}(x) \leq 1$

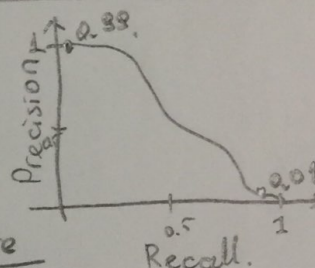
Suppose we want to predict  $y=1$  (cancer) only if very confident.

Predict 1 if  $h_{\theta}(x) \geq 0.7$  }  $\rightarrow$  Higher precision.  
 Predict 0 if  $h_{\theta}(x) < 0.7$  }  $\rightarrow$  Lower recall.

Suppose we want to avoid missing too many cases of cancer (avoid false negatives).

1 if  $h_{\theta}(x) \geq 0.3$  }  $\rightarrow$  Higher recall  
 0 if  $h_{\theta}(x) < 0.3$  }  $\rightarrow$  Lower precision.

More generally: Predict 1 if  $h_{\theta}(x) \geq \text{threshold}$



F<sub>1</sub> score (F score)

	Precision (P)	Recall (R)	Average	F <sub>1</sub> score
Algor. 1	0.5	0.4	0.45	0.444
Algor. 2	0.7	0.1	0.4	0.175
Algor. 3	0.02	1.0	0.51	0.0392

$$F_1 \text{ Score} = 2 \frac{PR}{P+R}$$

### ⑥. USING LARGE DATA SETS.

#### 6.1. Data for M.L.

Designing a high accuracy learning strn:

Large data rationale: Assume feature  $x \in \mathbb{R}^{n+1}$  has sufficient info to predict  $y$ .

Algorithms:

- Perceptron (Logistic Regression)
- Winnow
- Memory-based
- Naive Bayes.

$\Rightarrow$  Use a learn. alg. with many paramtrs (hidden units)  
 $\rightarrow$  low bias algorithms.  $J_{\text{train}}(\theta)$  will be small.

$\Rightarrow$  Use a very large training set  
 $\rightarrow$  low variance

$$J_{\text{train}}(\theta) \approx J_{\text{test}}(\theta) \text{ will be small}$$